



Reinforcement Learning from Human Feedback (RLHF)

x: "write me a poem about
the history of jazz"



preference data


maximum
likelihood



reinforcement learning

Direct Preference Optimization (DPO)

x: "write me a poem about
the history of jazz"



preference data

maximum
likelihood

