


(a) Latency of attention kernels.



(b) End-to-end latency with different block sizes.