

Thought: To diagnose the problem from the alert, I need to start by getting the current alerts.

Action: using tool

GetAlerts

Observation: High cost spend for vCPU and memory resources in the otel-demo namespace.

Thought: Received the current alerts. My next step is to gather more information about deployments...

Action: using tool
kubectl get deployment
adservice -n otel -o yaml

NL2Kubectl

Observation:

spec:

progressDeadlineSeconds: 600

replicas: 5 ...

Thought: The "adservice" deployment has many replicas. I should consider decreasing the resources allocated such as *decreasing the cpu and memory requests or deleting replicas*.

Action: using tool
kubectl patch deployment adservice
-n otel -p '{"spec": {"replica": 2}}'

NL2Kubectl

Observation: deployment.apps/adservice patched

The alert was not cleared after scaling the replicas.